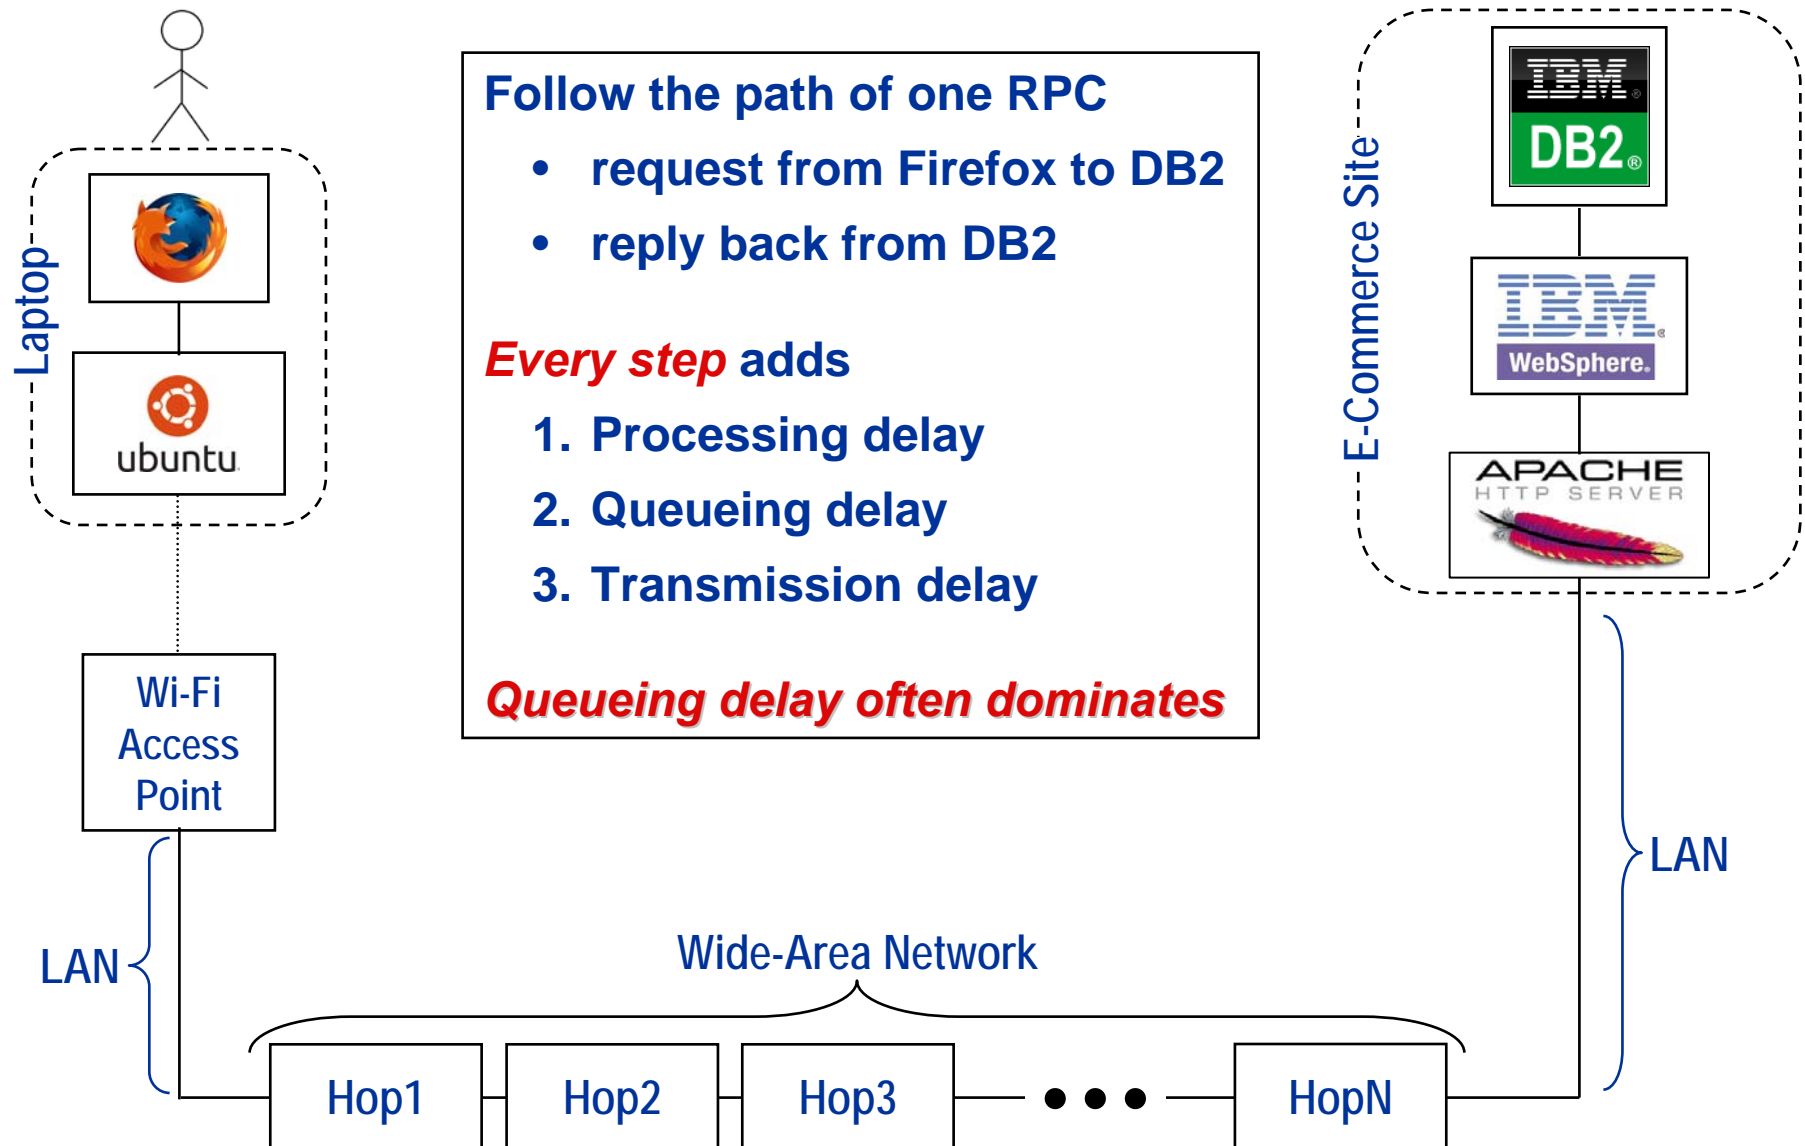


# Why Are Distributed Systems Slow?

# End to End RPC Latency



# Queueing Theory 101

- 1. Efficiency** (i.e. resource utilization: e.g. CPU, network, disk, etc.)
- 2. Crispness** (i.e., response time: user perception of quality)
- 3. Freedom** (no advance reservations; just use when needed)

**You can have at most 2 out of 3**

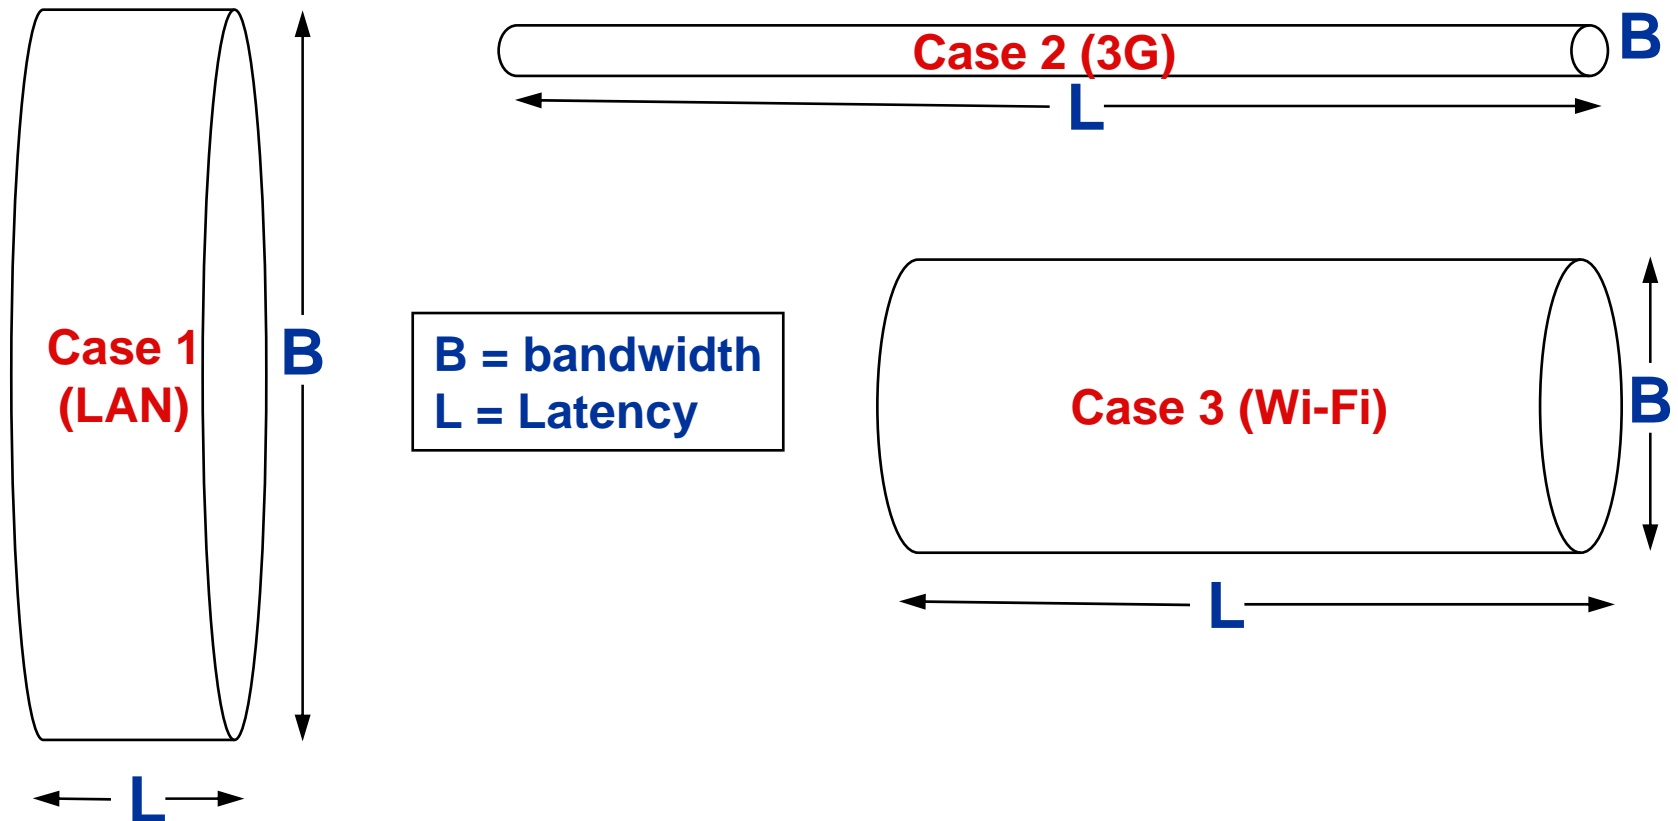
*you can't have all 3 – no amount of cleverness helps*

*best you can achieve is satisfactory tradeoff*

**Alas**

- bean counters demand high efficiency (i.e. high utilization)
- freedom is non-negotiable
- so, crispness falls victim → long queues and high queueing delays

# Latency and Bandwidth



Delay-Bandwidth product  $\approx$  max number of bits in flight

# Latency

*Latency is the killer, not bandwidth*

- fundamentally more difficult to improve

Bandwidth improved through parallelism

- fatter pipes, more lanes on highway, more checkout clerks at store, ...
- does cost money, but not fundamentally difficult

Latency is much, much, much harder to improve

- typically requires deep structural changes
- e.g shorten distance, increase max speed tolerated, reduce path length

Many software and systems trends increase latency

- increased use of layers (middleware, external libraries (DLLs), VMs, ...)
- security technologies (firewalls, overlay networks, ...)

# Impact of Long Latency

Synchronous model of RPC becomes infeasible

- *coast-to-coast US  $\approx 16$  ms at speed of light*
- round-trip RPC > 30 ms
- larger distances (e.g. trans-Pacific) will make matters worse

30 ms is a long time in terms of lost processing opportunity

- 3 million instructions on an early 1990s processor (100 MIPS )
- much higher on faster processors and multi-core machines
- a modern 3GHz single-core x86 processor is ~1000-1500 MIPS
- so 30 ms is ~10-15 million instructions (or more, with multicore)

*Can't afford to hide real-world asynchrony*

- RPC is a synchronous abstraction
- inadequate for many emerging use cases
- more complex, asynchronous models needed
- more difficult to program and get correct

# “Trust But Verify”

*aka Optimistic Methods*

Consider a very high latency distributed system

- e.g. Mars rover has ~8 minutes RTT for tele-operation  
(8m minimum, 48m maximum; depends on Earth-Mars distance)
- RPC model simply won't work; too slow and unresponsive
- only an optimistic method, with verification has hope (asynchronous)

*Information at one end about the other is always 8 min stale*

latency-imposed limit, in this case speed of light

Approach: include a predicate to validate before execution

- “When we last heard from you here was the situation...”
- “If this is still true, please do the following ...”
- “If not, just tell us what happened and we will try to give you guidance”

## Even more extreme:

- transatlantic communication before invention of telegraph  
age of sail
- one-way latency was ~1 week; RTT was ~2 weeks
- how was command and control done?  
e.g. King George and Lord Cornwallis during the American Revolution?

## For deeper discussion see optional reading:

*“Fundamental Challenges in Mobile Computing”*

Satyanarayanan, M.,

Proceedings of the Principles of Distributed Computing, 1996

## “Trust but Verify” is an optimistic approach

- contrast with simpler *pessimistic* approaches (lock or lease)
- more important as network latency dominates processing speed
- favors *liveness*, with controlled relaxation of *safety*