

Announcements

1. Problem Set 1 out today, due Thursday, Feb 12 (start of class)
2. Mid-term exam on Thursday, Feb 26 (in class)

3. Callback

Basic Idea

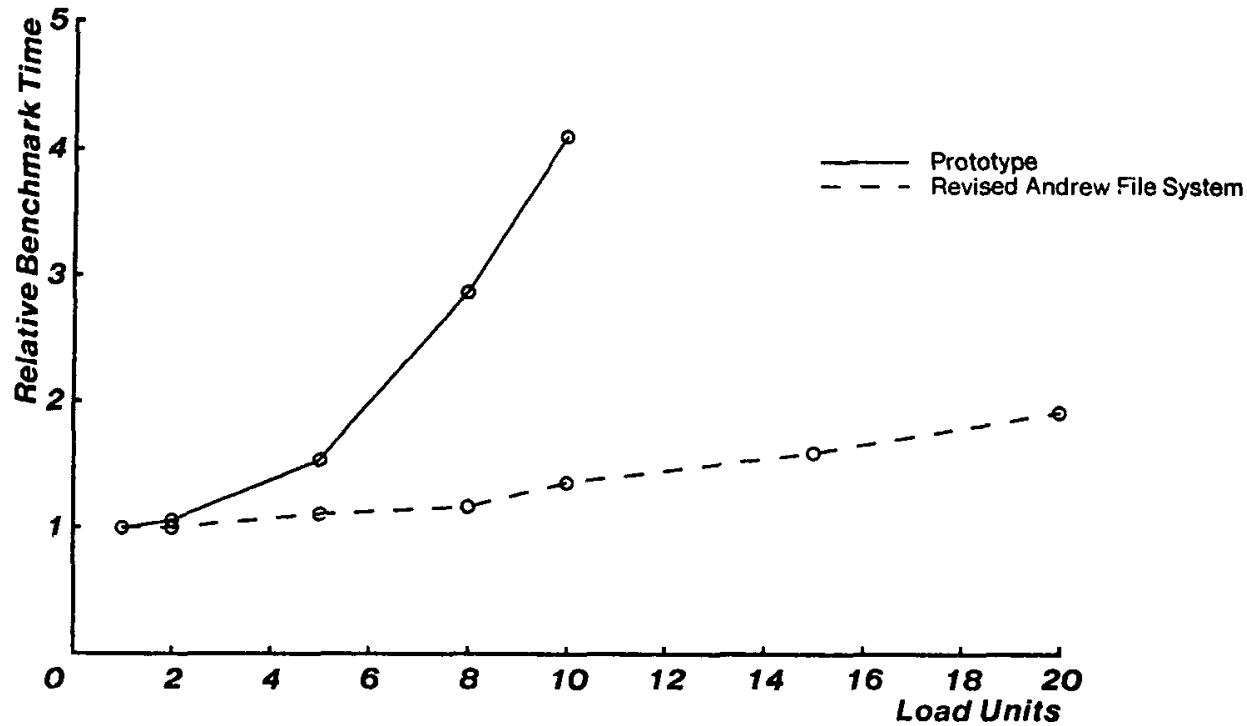
- *targeted notification of caching sites*
- master copy tracks sites with cached data
- typically done at coarse granularity (e.g. entire file)
can be made to work with byte ranges
- on update, all sites with cached data notified (*“callback”*)

Original Use

- AFS-2 (circa 1985)

Advantages

- excellent scalability for Unix workloads
- *zero network traffic for read of cached-valid objects*
- precursor to caching for disconnected operation
- biases read performance in favor of write-performance



Disadvantages

- sizable state on server
- complexity of tracking cached state on clients
- *silence ambiguous for client*
 - network failure → lost callbacks
 - periodic “keepalive” probes
 - data could be stale between probes
- NAT networks with masquerading firewalls

4. Leases

Basic Idea

Caching site obtains finite-duration control from master copy

duration is called “lease period”, typically few seconds

multiple sites can obtain read lease; only one can get write lease

- leases have to be renewed, else control expires with lease

Limiting cases

- lease duration = 0 → Check on Use
- lease duration = ∞ → callback (Targeted Notification)

Original Use

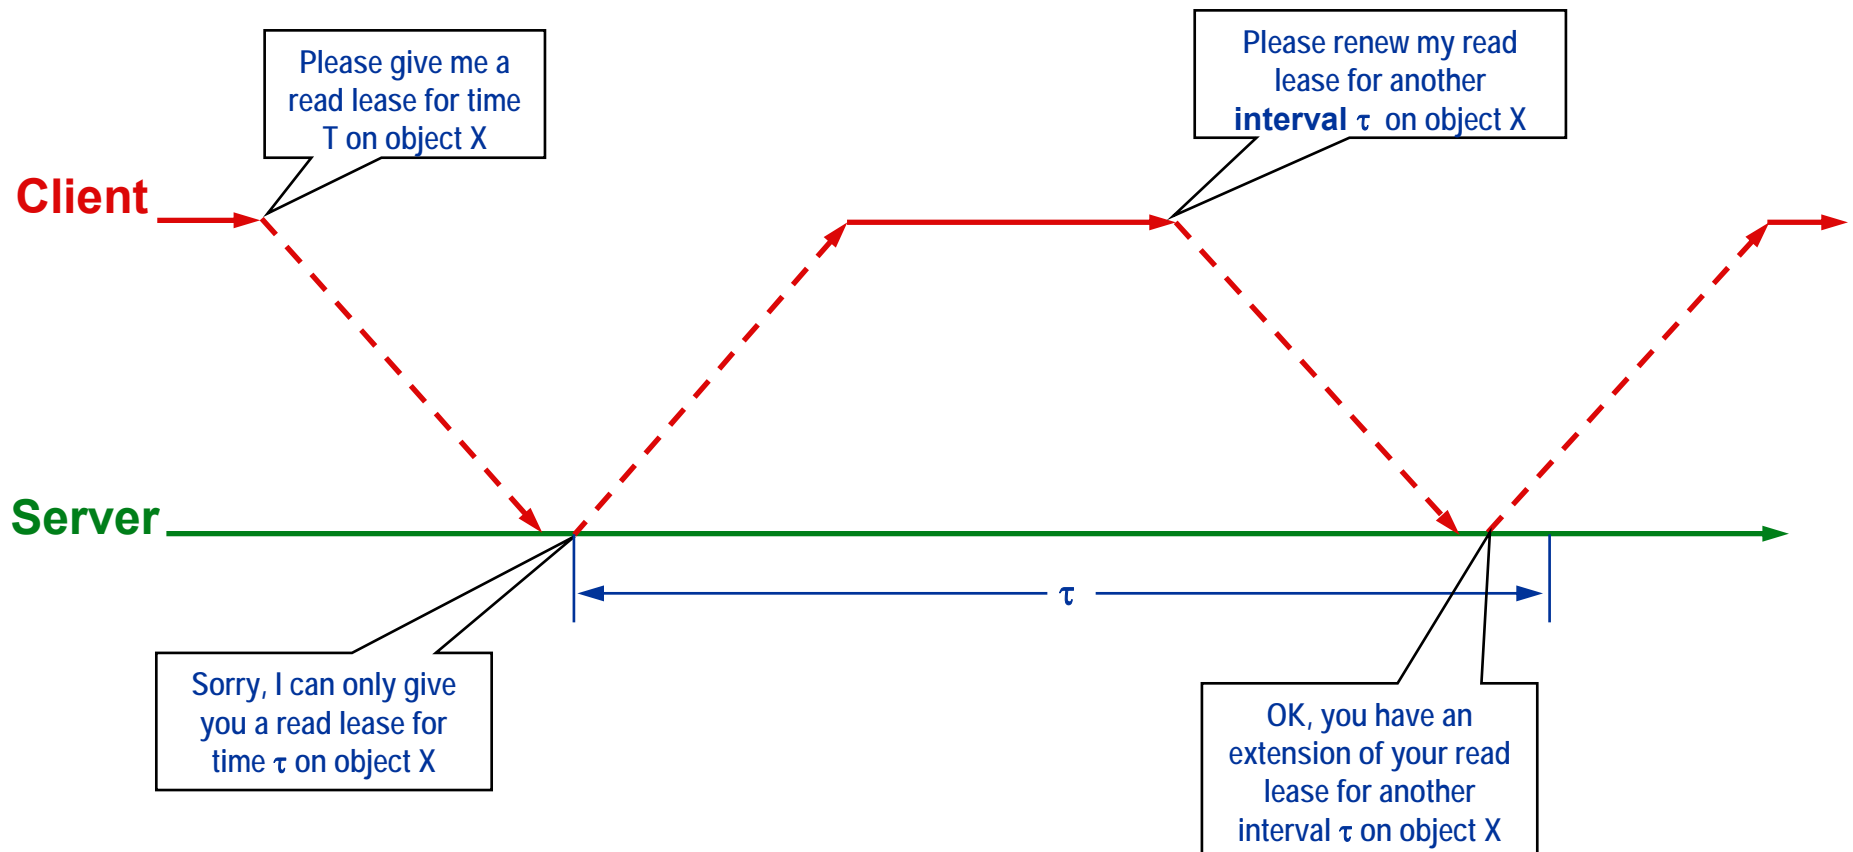
- V System (circa 1985) [Gray89]

Advantages

- generalizes the check on use and callback schemes
- lease duration can be tuned to adapt to mutation rate
Lease duration is a clean tuning knob for design flexibility
- conceptually simple yet flexible
- *time becomes a hidden communication channel*

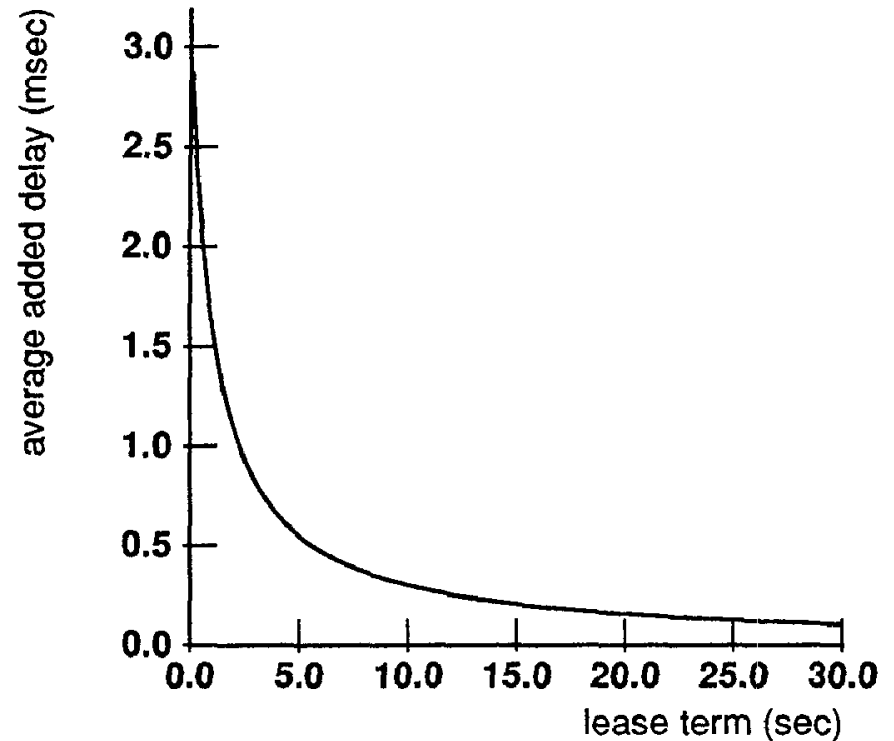
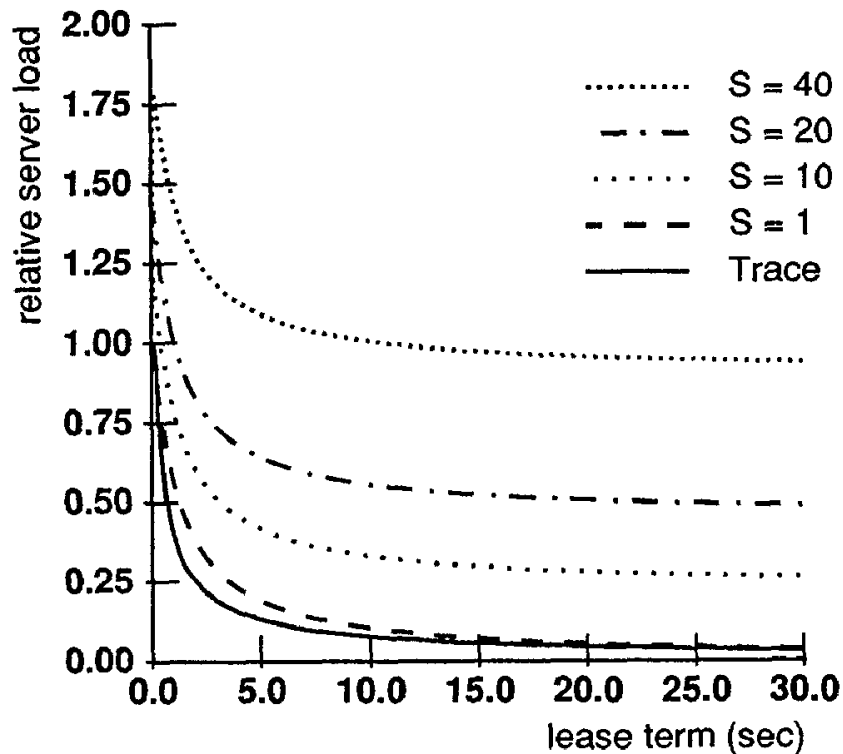
Lease Renewal

At most one write lease or multiple read leases at any one time



Disadvantages

- lease-holder has total autonomy during lease; revocation?
- writers delayed while read lease holders complete their leases
- more traffic than callback (but less than check on use)
 keepalives for callback only one per server, not per lease



5. Skip Scary Parts

Basic Idea

- *When write-sharing detected, turn off caching everywhere*
All references go directly master copy
- Resume caching when write-sharing ends

Original Use

- *Sprite (circa 1987)* (in conjunction with check on use)

Advantages

- Precise single-copy semantics (even at byte-level consistency)
- Excellent fallback position
Exemplifies good engineering: “Handle average case well; worst case safely”
Coda protocol achieves this differently: treats write-sharing as exception (conflict)
- Good adaptation of caching aggressiveness to workload

Disadvantages

- Server maintains state
- *Server aware of every use of data (open)*

6. Faith-Based Caching

Basic Idea

- *Blindly assume cached data is valid for a while*
- Periodically check (based on time since last check)
- No communication needed during trust period

Original use

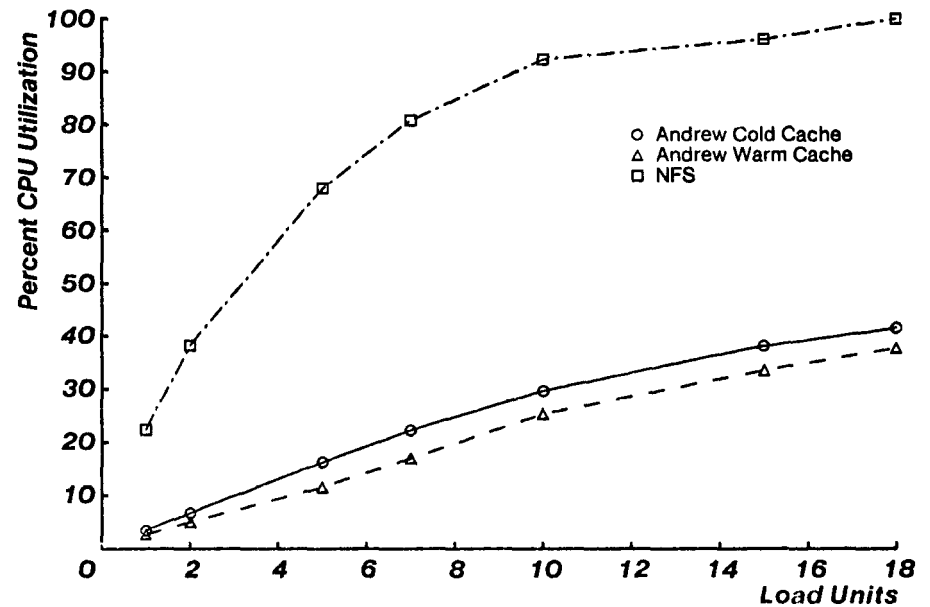
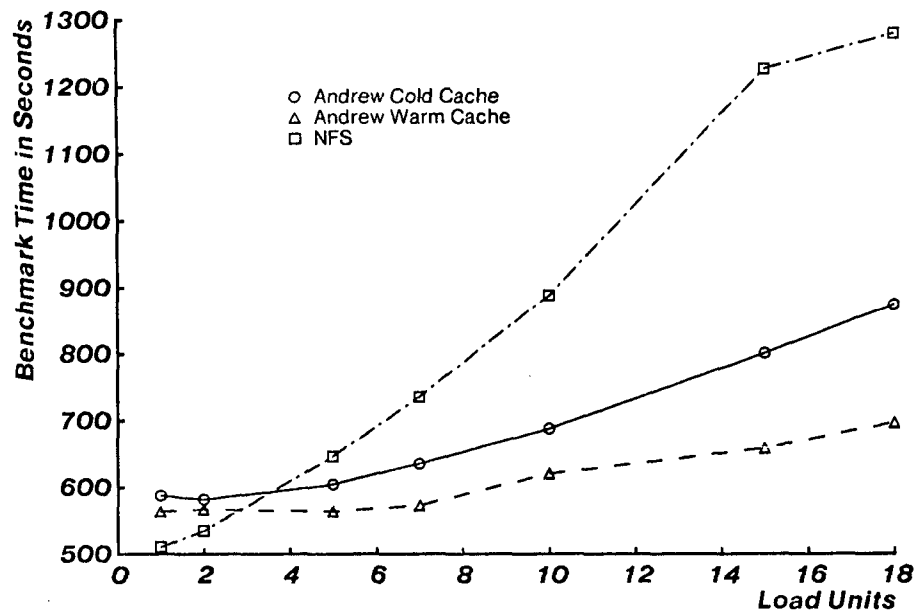
- *Sun NFSv3* file system
cached file blocks assumed current for X seconds
X = 3 for files, 30 for directories
- Small variant is a *TTL field* for each object
used in web caching, gives content creator modicum of control

Advantages

- Simple implementation
- *Server is stateless*

Disadvantages

- User-visible inconsistencies sometimes seen (make)
- Blind faith sometimes misplaced!
- Not as efficient as callback-based schemes



7. Pass the Buck

Basic Idea

- *Let the user trigger cache revalidation* (hit “reload”)
- Otherwise, all cached copies assumed valid forever
- Equivalent to infinite-TTL faith-based caching

Original use

- Arose in the context of the Web

Advantages


- Trivial to implement, no server changes
- Avoids frivolous cache maintenance traffic

Disadvantages

- Places burden on user
- User may be clueless about level of consistency needed
- Assumes existence of user: pain for write scripts/programs

Cache Consistency Strategies

1. *Broadcast invalidations*
2. *Check on Use*
3. *Callbacks*
4. *Leases*
5. *Skip Scary Parts*
6. *Faith-based Caching*
7. *Pass the Buck*



Many minor variants over the years, but these have withstood the test of time