

Concurrency Review 2

15-440/640 Spring 2015

February 10

Last Lecture

- Why we want concurrency
- Approaches to implementing concurrency
- Need to manage concurrency
(Synchronization)
- Mutex (lock) operations
- Lock granularity
- Deadlocks and avoidance

Beyond “simple” locks

- Semaphores
- Condition variables

Semaphores

- Non-negative counting variables
- Two atomic operations:
 - $P(s)$ – block until $s > 0$, then $s--$
 - $V(s)$ – $s++$
 - P also called wait, down, test
 - V also called post, up, increment
- Can be thought of as a “counting” lock

Semaphores vs. Mutexes

- Mutexes
 - Binary
 - Really intended for locks
 - Concept of “holding” lock – only holder can unlock
- Semaphores
 - Store value
 - Intended as a signal
 - Any thread can do P, V



Semaphore Uses


- Age old problem:



- Waiting for any one of N resources
- Use semaphore, initialized to N
- First N visitors calling P() will not have to wait
- $N+1^{\text{st}}$ visitor calling P() will block
- Will be released when any one visitor is done, calling V()

Semaphores Uses

- Often used with mutex
- E.g. Producer-consumer on n-element buffer



Semaphores for producer-consumer

- Use a mutex to coordinate access to buffer,
2 semaphores to manage number of resources

Semaphore free_slots, initialized to n

Semaphore filled_slots, initialized to 0

Mutex L

Buffer[n]

Producers:

P(free_slots)

Lock(L)

Add item to Buffer

Unlock(L)

V(filled_slots)

Consumers:

P(filled_slots)

Lock(L)

Remove item from Buffer

Unlock(L)

V(free_slots)

- Ok, so far, we have dealt with synchronizing access to shared data or resources
- Sometimes, we also want to coordinate when threads run

Event ordering between threads

- How can we make sure thread A reaches point X before B reaches Y ?
- One solution: use a semaphore, initialized to 0

Thread A

...

X: $V(s)$

...

Thread B

...

P(s)

Y: ...

DIY scheduler

- Non-preemptive time sliced execution
- Each thread marked with “yield” points:

```
void yield() {  
    V(s_sched);  
    P(s[thread_id]);  
}
```

- Scheduling thread main loop:

```
while(1) {  
    P(s_sched);  
    j = get_next_thread_id_to_run();  
    V(s[j]);  
}
```

- Initial value of s_sched determines number of concurrent threads.

On your mark, get set, go!

- Suppose several threads want to wait for some event to occur, then all resume
- How can we do this?
 - Maybe use a semaphore. N threads blocked on it. When event occurs, do $V()$ N times.
 - But last part is not atomic, can run into problems

Condition Variables

- Three operations:
 - Wait() – wait for condition
 - Signal() – wake one waiting thread
 - Broadcast() – wake all waiting threads
- Condition variables must be used with a mutex (we'll come back to this)

Barrier Synchronization

- Barrier: a point in the code where all threads stop, and wait for all others to get there too
- Common paradigm:
 - Threads perform computation phase A; wait for all to finish, then start phase B
 - Useful in many simulations:
 - Threads simulate 1 “time step” for their slice of the world
 - Wait for all threads, then exchange state
 - Repeat

Implementing a Barrier

- We can use a CV to wake everyone waiting
- When do we trigger the broadcast to the CV?
 - When the last thread reaches the barrier!
 - Use a counter to track number of waiting threads
- Pseudocode:

```
mutex L;  
int num_waiting=0;  
condvar cv;
```

```
barrier() {  
    lock(L);  
    num_waiting++;  
    if (num_waiting==NUM_THREADS) {  
        num_waiting=0;  
        broadcast(cv);  
        unlock(L);  
    } else {  
        unlock(L);  
        wait(cv);  
    }  
}
```

What if some thread
is still here when
broadcast is called?

- Not quite right – there is a *race condition*

Race Conditions

- A race condition exists if a thread needs to reach point x before another thread reaches point y to make the program work
 - In example barrier code, if one thread does not reach wait by time broadcast is called, it won't get past the barrier
 - Furthermore, at next barrier, num_waiting won't reach NUM_THREADS, so no one calls broadcast and every thread is left blocked!

Condition Variables to the rescue

- CVs are designed to fix this particular race
- Wait() must be used with a mutex:

Lock(L)

Wait(CV,L) – atomically unlock L and wait
when resumed, will have L locked

Unlock(L)

Fixed Barrier

- Barrier pseudocode with race removed:

```
mutex L;
int num_waiting=0;
condvar cv;

barrier() {
    lock(L);
    num_waiting++;
    if (num_waiting==NUM_THREADS) {
        num_waiting=0;
        broadcast(cv);
    } else wait(cv, L);
    unlock(L);
}
```

Summary of semaphores, CVs

- Semaphores
 - Signaling methods with an associated count
 - Great for coordinating access to a set of resources
 - Can also coordinate when threads run
- Condition Variables
 - Used to signal an event to multiple waiting threads
 - Great for barrier synchronization

What about other languages?

- We have illustrated concurrency concepts in C, since everything is explicit
- Most languages provide similar primitives
- Language syntax may hide synchronization operations, but they are really implemented in the same way

Java Threads

- Provide a class that implements Runnable:

```
public class MyClass implements Runnable {  
    public void run() {  
        /* do stuff */  
    }  
}
```

- To launch thread, instantiate a Thread object:

```
Thread t = new Thread( new MyClass() );
```

- Exits when run() returns
- Join operation:

```
t.join();
```

Java Threads - alternative

- Can also subclass Thread:

```
public class MyClass extends Thread {  
    public void run() {  
        /* do stuff */  
    }  
}
```

- Launch thread by creating instance and calling start:

```
( new MyClass() ).start();
```

Java Concurrency Management

- Java provides “synchronized” methods:

```
public class Account {  
    private float balance;  
    public synchronized void deposit(float v) {  
        balance+=v;  
    }  
}
```

- Ensures mutual exclusion among all synchronized methods of an object instance

Synchronized, behind the scenes

- All Java objects have an “intrinsic lock”
- Synchronized methods simply acquire the lock at start, and release at return

```
public class Account {  
    // implicit private mutex lock  
    private float balance;  
    public synchronized void deposit(float v) {  
        // implicit lock(lock)  
        balance+=v;  
        // implicit unlock(lock)  
    }  
}
```

- Intrinsic locks are reentrant

Synchronized, not magical

- Fine grained locking at object instance level

```
public class Account {  
    private float balance;  
    public synchronized void deposit(float v) {  
        balance+=v;  
    }  
    public synchronized void transfer(Account dest, float v) {  
        balance-=v;  
        dest.deposit(v);  
    }  
}
```

- Can still deadlock!
- Language construct is convenient, but may hide that fact that multiple locks are being acquired
- Non-synchronized methods can still mess things up

More control over locking

- Synchronized Statements:

```
synchronized( object ) {  
    / * do critical section */  
}
```

- Can use any object's intrinsic lock
- Can instantiate generic Object to use as a lock:
 private Object myLock = new Object();
- We can use this to control when locks are held, order of acquisition, etc.

Built-in Condition Variable Functionality

- wait, notify, notifyAll = wait, signal, broadcast
- Can use to implement barrier synchronization
- Can use to make a semaphore:

```
public class Semaphore {  
    private int count = 0;  
    public synchronized V() {  
        count++;  
        this.notify();  
    }  
    public synchronized P() {  
        while (count==0) {  
            try { this.wait() }  
            catch (InterruptedException e) {}  
        }  
        count--;  
    }  
}
```

Java Concurrency Summary

- Use Thread objects to launch new threads
- Synchronized methods and synchronized statements provide mutual exclusion
- All objects have intrinsic locks
- wait, notify, and notifyAll provide CV-like functionality in all objects

Concurrency Review

- Why we want concurrency
- Approaches to using concurrency
- Synchronizing primitives: Mutex, semaphore, condition variables
- Perils of concurrency: race conditions, deadlocks, performance issues
- C and Java examples